



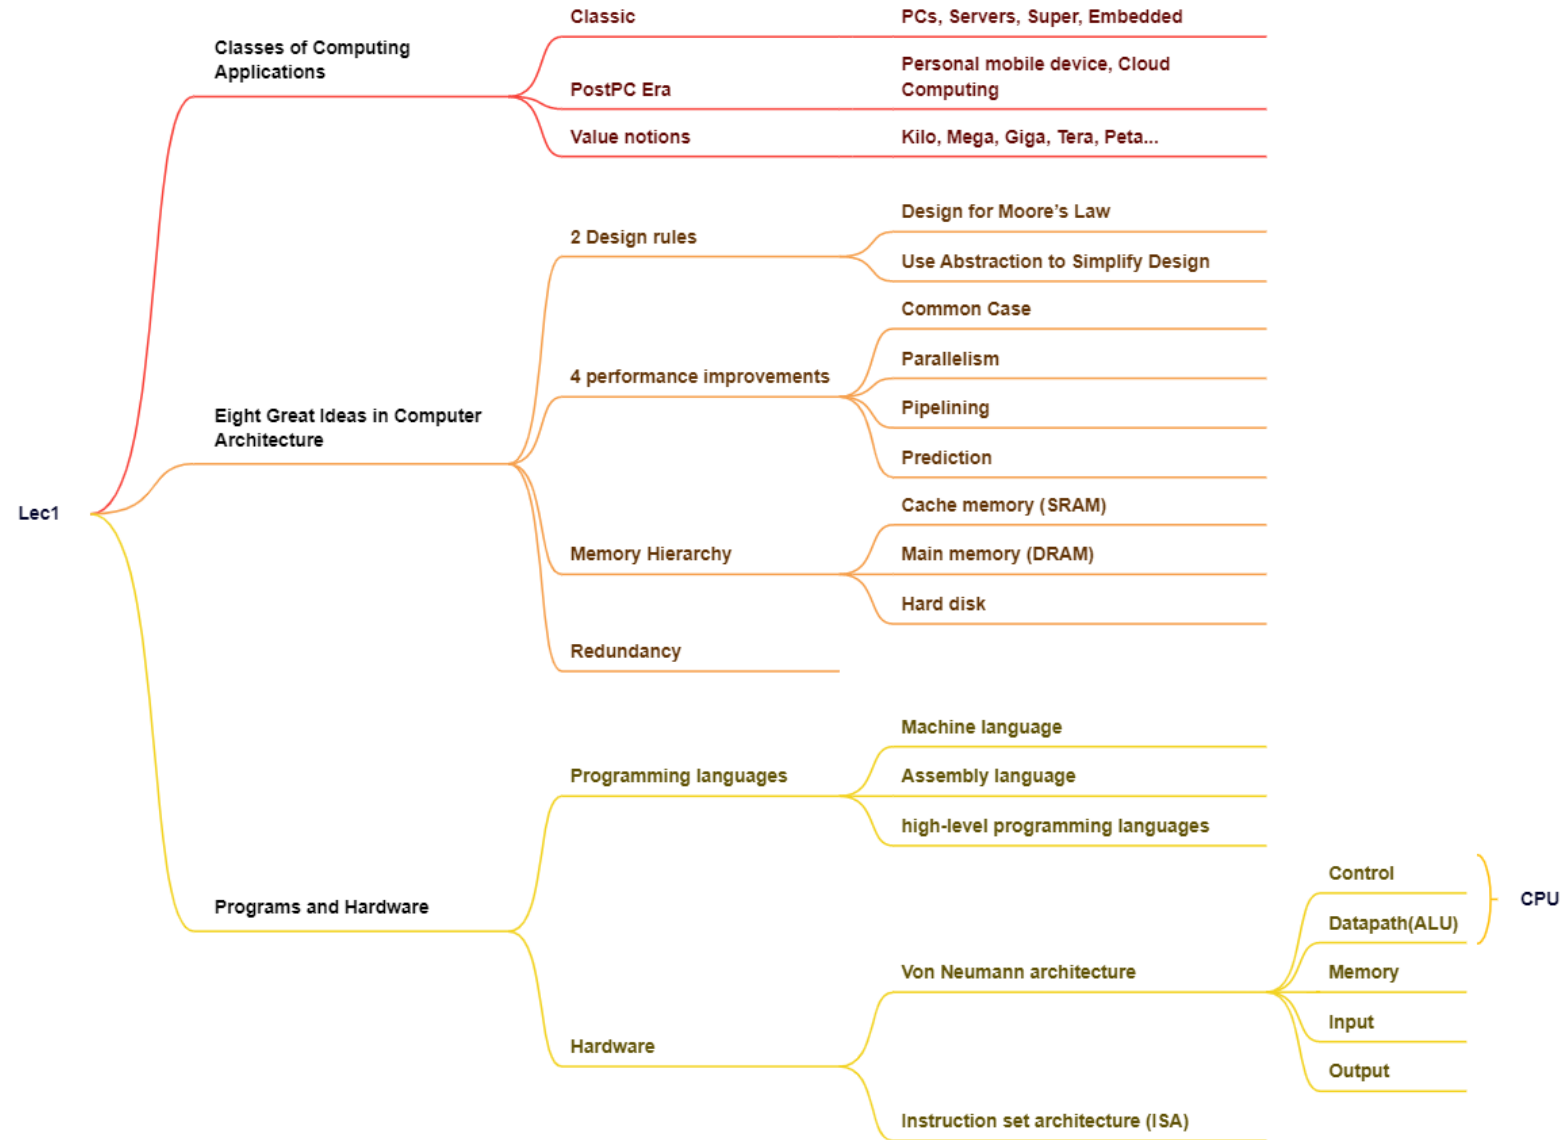
COMPUTER ORGANIZATION

Lecture 2 RISC-V Introduction

2025 Spring


This PowerPoint is for internal use only at Southern University of Science and Technology. Please do not repost it on other platforms without permission from the instructor.

Recap



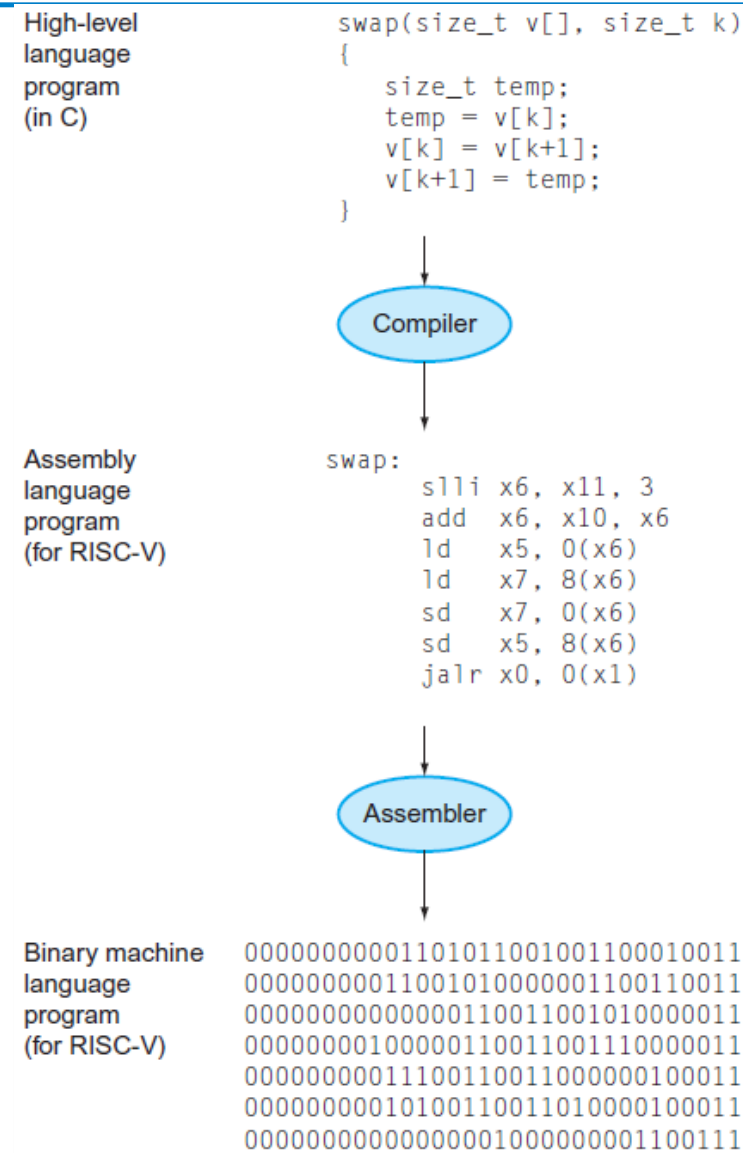
The Concept of a Computer

- Application software
 - Written in high-level language
- System software
 - Compiler: translates HLL code to machine code
- Operating System:
 - Handling input/output
 - Managing memory and storage
 - Scheduling tasks & sharing resources
- Hardware
 - Processor, memory, I/O controllers



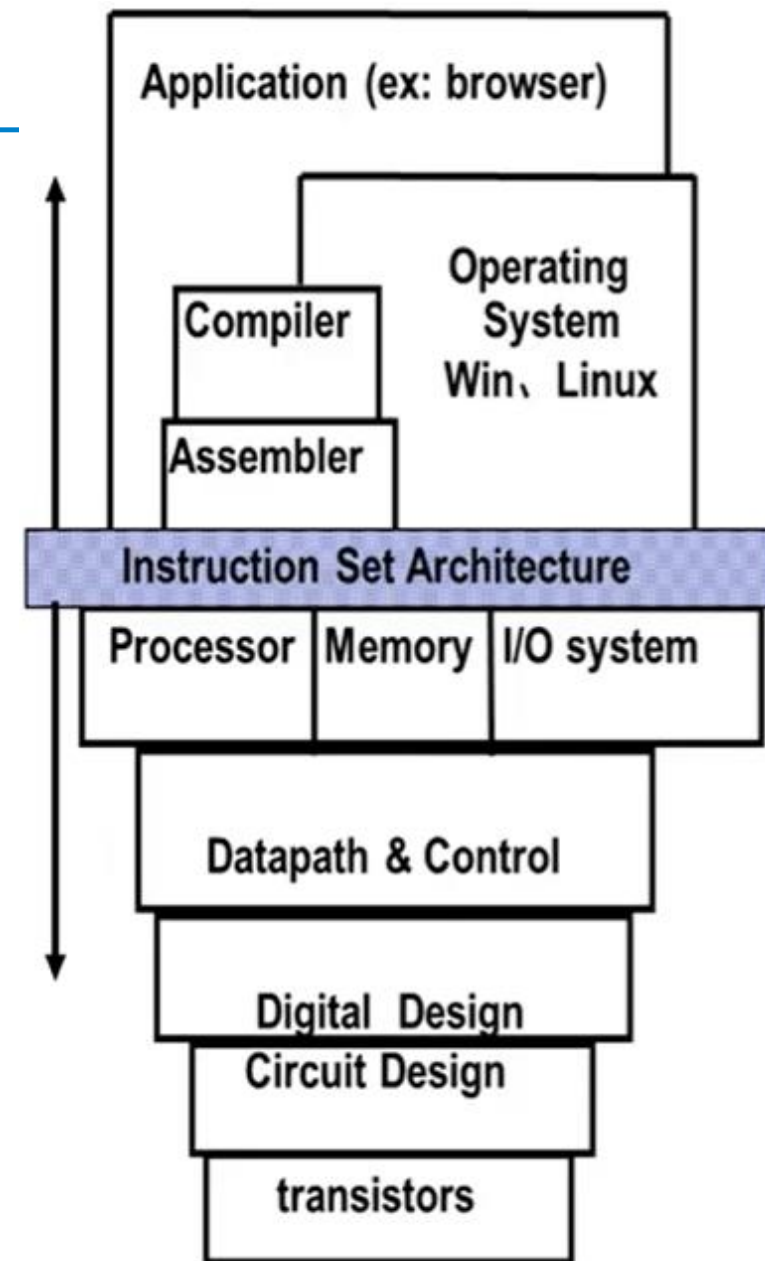
Levels of Program Code

- C program compiled into assembly language and then assembled into binary machine language.
- High-level language
 - Level of abstraction closer to problem domain
 - Provides for productivity and portability
- Assembly language
 - Textual representation of instructions
- Machine language
 - Hardware representation
 - Binary digits (bits)
 - Encoded instructions and data



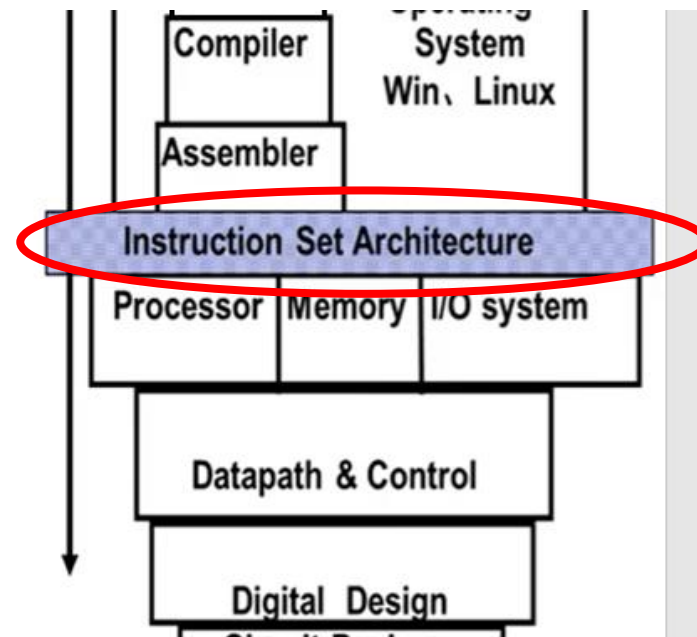
Abstractions

- Abstraction helps us deal with complexity
 - Hides lower-level details
- Instruction Set Architecture (ISA) or Computer Architecture
 - The hardware/software interface
 - Includes instructions, registers, memory access, I/O, and so on
- Operating system hides details of doing I/O, allocating memory from programmers



Instruction Set Architecture (ISA)

- A set of assembly language instructions (ISA) provides a link between software and hardware.
- Given an instruction set, software programmers and hardware engineers work more or less independently.
- Common types of ISA: RISC, CISC
- Examples:
 - IBM370/X86 (CISC)
 - **RISC-V** (RISC)
 - MIPS (RISC)
 - ARM (RISC)



Instruction Set Architecture (ISA)

- Instructions: CPU's primitive operations
 - Instructions performed one after another in sequence
 - Each instruction does a small amount of work (a tiny part of a larger program).
 - Each instruction has an operation applied to operands,
 - and might be used to change the sequence of instructions.
- CPUs belong to “families,” each implementing its own set of instructions
- CPU's particular set of instructions implements an Instruction Set Architecture (ISA)
 - Examples: ARM, Intel x86, MIPS, RISC-V, PowerPC...

Instruction Set Architecture

- CISC
 - Complex Instruction Set Computer
 - Variable instruction length
 - Much more powerful instructions
 - Hardware intensive instructions (more transistors)
 - e.g. x86
- RISC
 - Reduced Instruction Set Computer
 - Fixed instruction size
 - Simple instructions (load/store)
 - Emphasizes more on software (compiler)
 - e.g. MIPS, ARM, PowerPC, RISC-V

Which ISAs “win”

- The big winners: x86/x64 (servers) and Arm (phones/embedded)
 - Neither are the cheapest nor the best architectures available...
 - They won because of the legacy ecosystem
- But since our focus is understanding how computers work, we choose learning RISC-V
- Learn to program in assembly language, e.g. RISC-V
 - Best way to understand what compilers do to generate machine code
 - Best way to understand what the CPU hardware does

And the Road To Future Classes

- CS302 Operation Systems
 - OS needs a small amount of assembly for doing things the "high level" language doesn't support
 - Such as accessing special resources
- CS323 Compilers
 - Learn how to build compilers. A compiler goes from source code to assembly language.
- CS301 Embedded System
 - Assembly or a combination of high-level languages and inline assembly are commonly used to achieve efficient execution in resource-constrained environments.
- CS315 Computer Security
 - Exploit code ("shell code") is often in assembly and exploitation often requires understanding the assembly language & calling-convention of the target



What is RISC-V

- Fifth generation of RISC design from UC Berkeley
- A high-quality, license-free, royalty-free RISC ISA specification
 - Implementers do not pay any royalties
 - Large community of users riscv.org: industry, academia
 - Full software stack
- Appropriate for all levels of computing system, from microcontrollers to supercomputers
 - 32-bit, 64-bit, and 128-bit variants (**we're using 32-bit(RV32) in lectures and labs, textbook uses 64-bit**)
- Standard maintained by non-profit RISC-V Foundation

A Basic Assembly Instruction

- C code: `a = b + c ;`
- Assembly code: (human-friendly machine instructions)
`add a, b, c # a is the sum of b and c`
- Machine code: (hardware-friendly machine instructions)
`0000 0000 1100 0101 1000 0101 0011 0011`
- Translate the following C code into assembly code:
`a = b + c + d + e;`

<code>add a, b, c</code>		<code>add a, b, c</code>
<code>add a, a, d</code>	or	<code>add f, d, e</code>
<code>add a, a, e</code>		<code>add a, a, f</code>
- Instructions are simple: fixed number of operands (unlike C)
- A single line of C code is converted into multiple lines of assembly code
- Some sequences are better than others... the second sequence needs one more (temporary) variable `f`

A Basic Assembly Instruction

- In Previous example
 - `add a, b, c`
 - All arithmetic operations have this form
- *Design Principle 1: Simplicity favors regularity*
 - Regularity makes implementation simpler
 - Simplicity enables higher performance at lower cost
- Example
 - C code: `f = (g + h) - (i + j);`
 - Assembly code:

<code>add</code>	<code>t0,</code>	<code>g,</code>	<code>h</code>	<code># temp t0 = g + h</code>
<code>add</code>	<code>t1,</code>	<code>i,</code>	<code>j</code>	<code># temp t1 = i + j</code>
<code>sub</code>	<code>f,</code>	<code>t0,</code>	<code>t1</code>	<code># f = t0 - t1</code>

Assembly Variables: Registers

- Unlike HLL like C or Java, assembly does not have variables
- Assembly language operands are objects called registers
 - Limited number of special places to hold values, built directly into the hardware
 - 32 registers in RISC-V
 - Each RISC-V register is 32 bits wide called a “word” (RV32 variant of RISC-V ISA)
 - Registers have **no type**
 - Operation determines how register contents are interpreted
- *Design Principle 2: Smaller is faster*
 - registers: 32
 - main memory: millions of locations

RISC-V Registers


- x0 is special, always holds the value zero and can't be changed

Register	Alternative name	Description
x0	zero	the constant value 0
x1	ra	Return address
x2	sp	Stack pointer
x3	gp	Global pointer
x4	tp	Thread pointer
x5 – x7	t0 – t2	Temporaries
x8	s0/fp	Saved register/Frame pointer
x9	s1	Saved register
x10-17	a0-7	Function arguments/Return values
x18-27	s2-11	Saved registers
x28-31	t3-6	Temporaries

RISC-V Instructions

- Instructions have an opcode and operands

• E.a... add x1. x2. x3 # x1 = x2 + x3



- Instructions are fixed, 32bit long (machine code)
 - Note: Conversions between assembly to corresponding machine code will be taught in future lecture
- Each instruction uses one of these predefined formats:

Name (Field Size)	7 bits	5 bits	5 bits	3 bits	5 bits	7 bits	Comments
R-type	funct7	rs2	rs1	funct3	rd	opcode	Arithmetic instruction format
I-type	immediate[11:0]		rs1	funct3	rd	opcode	Loads & immediate arithmetic
S-type	immed[11:5]	rs2	rs1	funct3	immed[4:0]	opcode	Stores
SB-type	immed[12,10:5]	rs2	rs1	funct3	immed[4:1,11]	opcode	Conditional branch format
UJ-type	immediate[20,10:1,11,19:12]				rd	opcode	Unconditional jump format
U-type	immediate[31:12]				rd	opcode	Upper immediate format

Arithmetic Operations

- Addition:
 - Example: `add x1, x2, x3` (in RISC-V)
 - Equivalent to: $a = b + c$ (in C), where a, b, c in x1, x2, x3
- Subtraction
 - Example: `sub x3, x4, x5` (in RISC-V)
 - Equivalent to: $d = e - f$ (in C), where d, e, f in x3, x4, x5
- Example: how to do the following C statement?
`f = (g + h) - (i + j);`
 - f, ..., j in x19, x20, ..., x23
 - Break into multiple instructions:

```
add x5, x20, x21    # temp t0 = g + h
add x6, x22, x23    # temp t1 = i + j
sub x19, x5, x6      # f = t0 - t1
```

Register x0

- Very useful: always holds zero and can never be changed (does not require initialization)
- Ex: Copy a value from one register to another:
`add x3,x4,x0` (in RISC-V)
same as
`f = g` (in C)
- Or, whenever a value is produced and we want to throw it away, write to x0:
- By convention RISC-V has a specific no-op instruction
`add x0,x0,x0`
- Also, we will see x0 used later with “jump-and-link” instruction

Immediates

- Immediates are used to provide numerical constants
- *Design Principle 3: Make the common case fast*
 - Small constants are common
 - Immediate operand avoids loading from memory
- Syntax similar to add instruction, except that last argument is a number instead of a register
- Example: Add Immediate:

`addi x3,x4,10`

same as

`f = g + 10 (in C)`

`addi x3,x4,0`

same as

`f = g (in C)`

- No subtract immediate instruction, why?

Recall: Numeric Representations

- Decimal 35_{10} or 35_{ten}
- Binary 00100011_2 or 00100011_{two}
- Hexadecimal $0x23$ or 23_{hex}
0-15 (decimal) \rightarrow 0-9, a-f (hex)

Recall: Numeric Representations

- Unsigned Binary Integers
- Given an n-bit number

$$x = x_{n-1}2^{n-1} + x_{n-2}2^{n-2} + \cdots + x_12^1 + x_02^0$$

- Range: 0 to $+2^n - 1$
- Example
 - 0000 0000 0000 0000 0000 0000 0000 1011₂
= $0 + \dots + 1 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 1 \times 2^0$
= $0 + \dots + 8 + 0 + 2 + 1 = 11_{10}$
- Using 32 bits
 - 0 to +4,294,967,295

Signed Numeric Representations

- 2s-Complement Signed Integers
- Given an n-bit number

$$X = -x_{n-1}2^{n-1} + x_{n-2}2^{n-2} + \dots + x_12^1 + x_02^0$$

- Range: -2^{n-1} to $+2^{n-1} - 1$
- Example
 - $1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1100_2$
 $= -1 \times 2^{31} + 1 \times 2^{30} + \dots + 1 \times 2^2 + 0 \times 2^1 + 0 \times 2^0$
 $= -2,147,483,648 + 2,147,483,644 = -4_{10}$
- Using 32 bits
 - $-2,147,483,648$ to $+2,147,483,647$

Signed Numeric Representations

- Signed Negation
- Complement and add 1
 - Complement means $1 \rightarrow 0, 0 \rightarrow 1$


$$x + \bar{x} = 1111 \dots 111_2 = -1$$

$$\bar{x} + 1 = -x$$

- Example: negate +2
 - $+2 = 0000 \ 0000 \dots 0010_2$
 - $-2 = 1111 \ 1111 \dots 1101_2 + 1$
 $= 1111 \ 1111 \dots 1110_2$

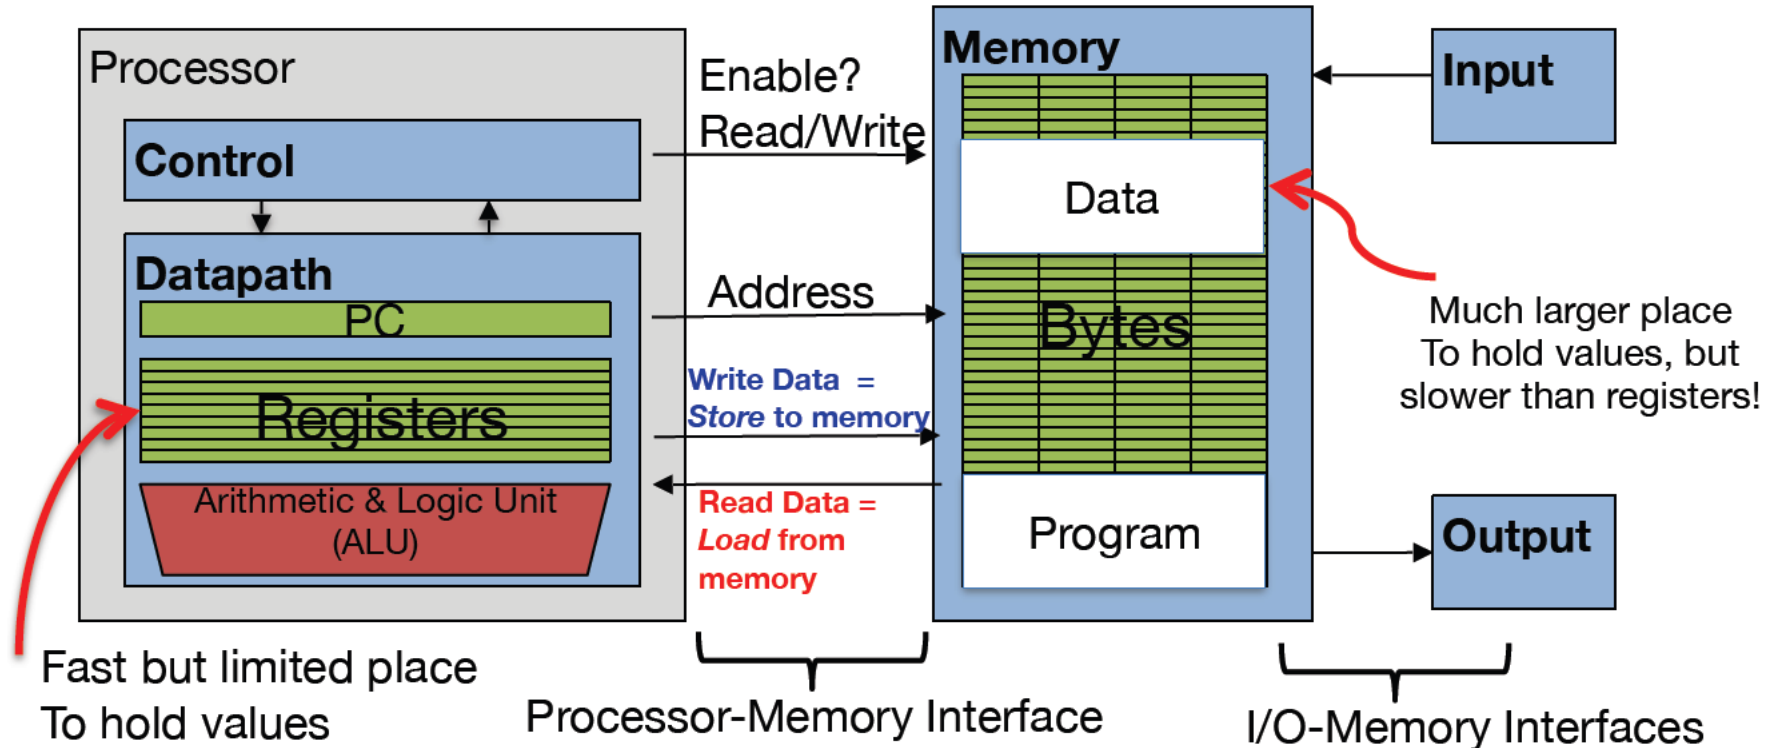
Immediates & Sign Extension

- Immediates are necessarily small
 - An I-type instruction can only have 12 bits of immediate (We'll see more details in future lecture)
- In RISC-V immediates are "sign extended"
 - So the upper bits are the same as the top bit
- Examples: 8-bit to 16-bit
 - +2: 0000 0010 => 0000 0000 0000 0010
 - -2: 1111 1110 => 1111 1111 1111 1110
- So for a 12bit immediate...
 - Bits [31:12] get the same value as Bit 11



Data Transfer Operations

- Registers vs. Memory
 - Arithmetic operations can only be performed on registers
 - Thus, the only memory actions are loads & stores

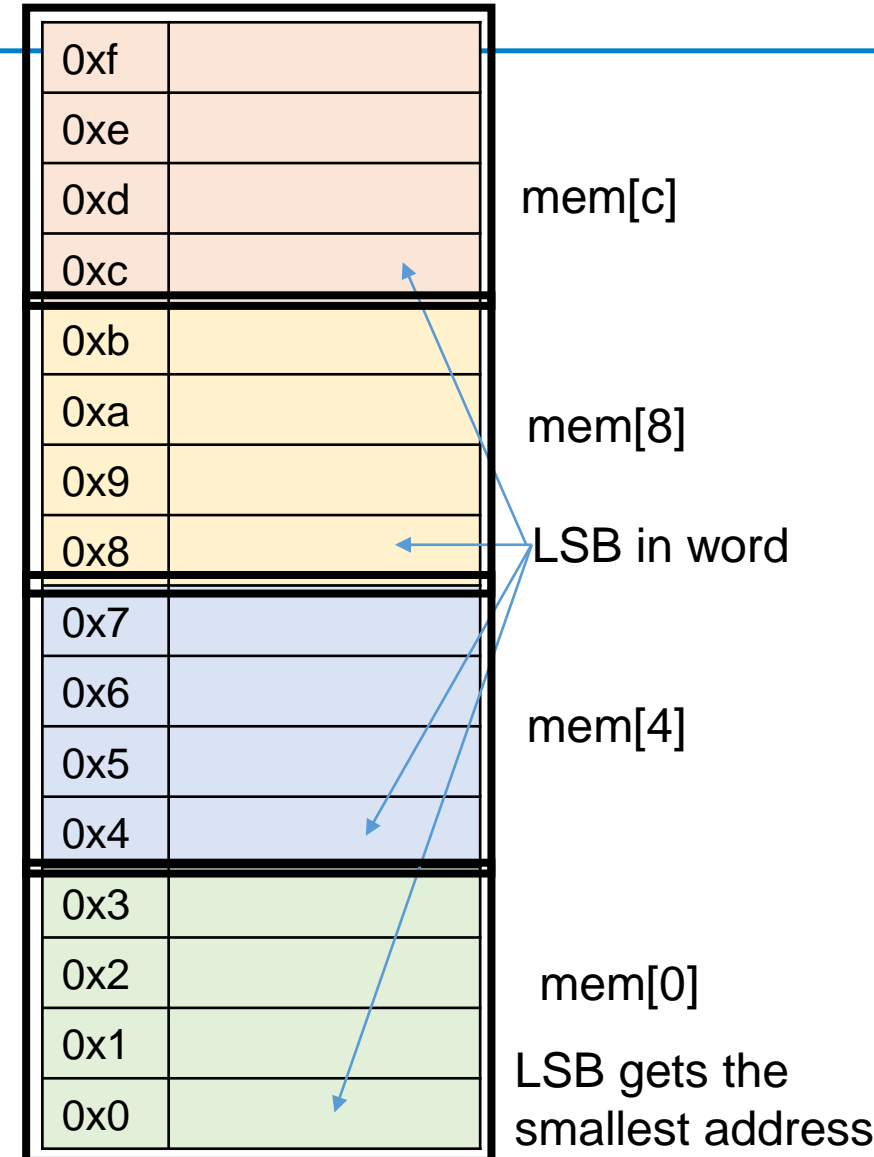


Speed of Registers vs. Memory

- Registers vs. Memory
 - Arithmetic operations can only be performed on registers
 - Thus, the only memory actions are loads & stores
- Given that
 - Registers: 32 words (128 Bytes)
 - Memory (DRAM): Billions of bytes (2 GB to 16 GB on laptop)
- How much faster are registers than DRAM??
- About 100-500 times faster!
 - in terms of latency of one access

Memory Addresses are in Bytes

- Data typically smaller than 32 bits, but rarely smaller than 8 bits (e.g., char type)
 - So everything is a multiple of 8 bits
- Remember, size of word is 4 bytes
- Memory is addressable to individual bytes
- Word addresses are 4 bytes apart
 - words take on the address of their least-significant byte
 - remember to keep words aligned
- RISC-V does not require words to be aligned in memory
 - But it is very **very bad !!!**
 - So in **practice**, RISC-V requires integers word-aligned !!!



Transfer from Memory to Register

- C code

```
int A[100];  
g = h + A[8];
```

- Assume: **x13** holds base register (pointer to A[0])
- Note: **32** is offset in bytes
- Offset must be a constant known at assembly time

- Using Load Word (**lw**) in RISC-V:

```
lw  x10, 32(x13)    # reg x10 gets A[8]  
add x11, x12, x10    # g = h + A[8]
```

...
base+32	A[8]	103
...
base+4	A[1]	830
base addr	A[0]	15

Transfer from Register to Memory

- C code

```
int A[100];  
A[10] = h + A[8];
```

- Assume: x13 holds base register (pointer to A[0])
- Note: 32, 40 is offset in bytes
- Offset must be a constant known at assembly time

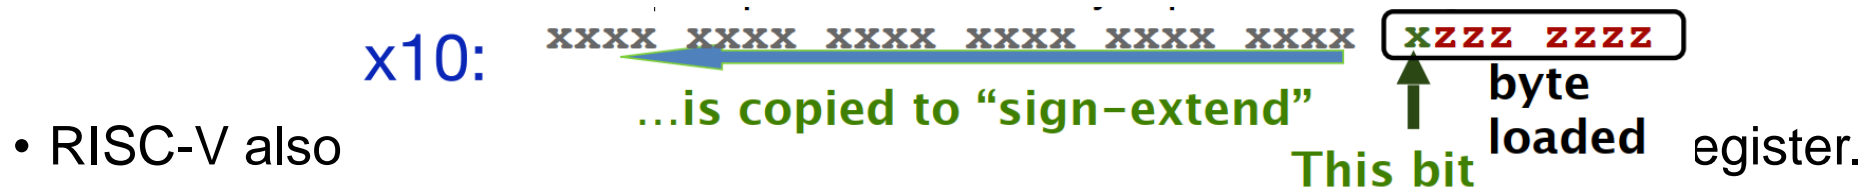
- Using Store Word (**sw**) in RISC-V:

```
lw  x10, 32(x13)      # reg x10 gets A[8]  
add x11, x12, x10      # g = h + A[8]  
sw  x11, 40(x13)      # A[10] = h + A[8]
```

- x13+32 and x13+40 must be multiples of 4 to maintain alignment

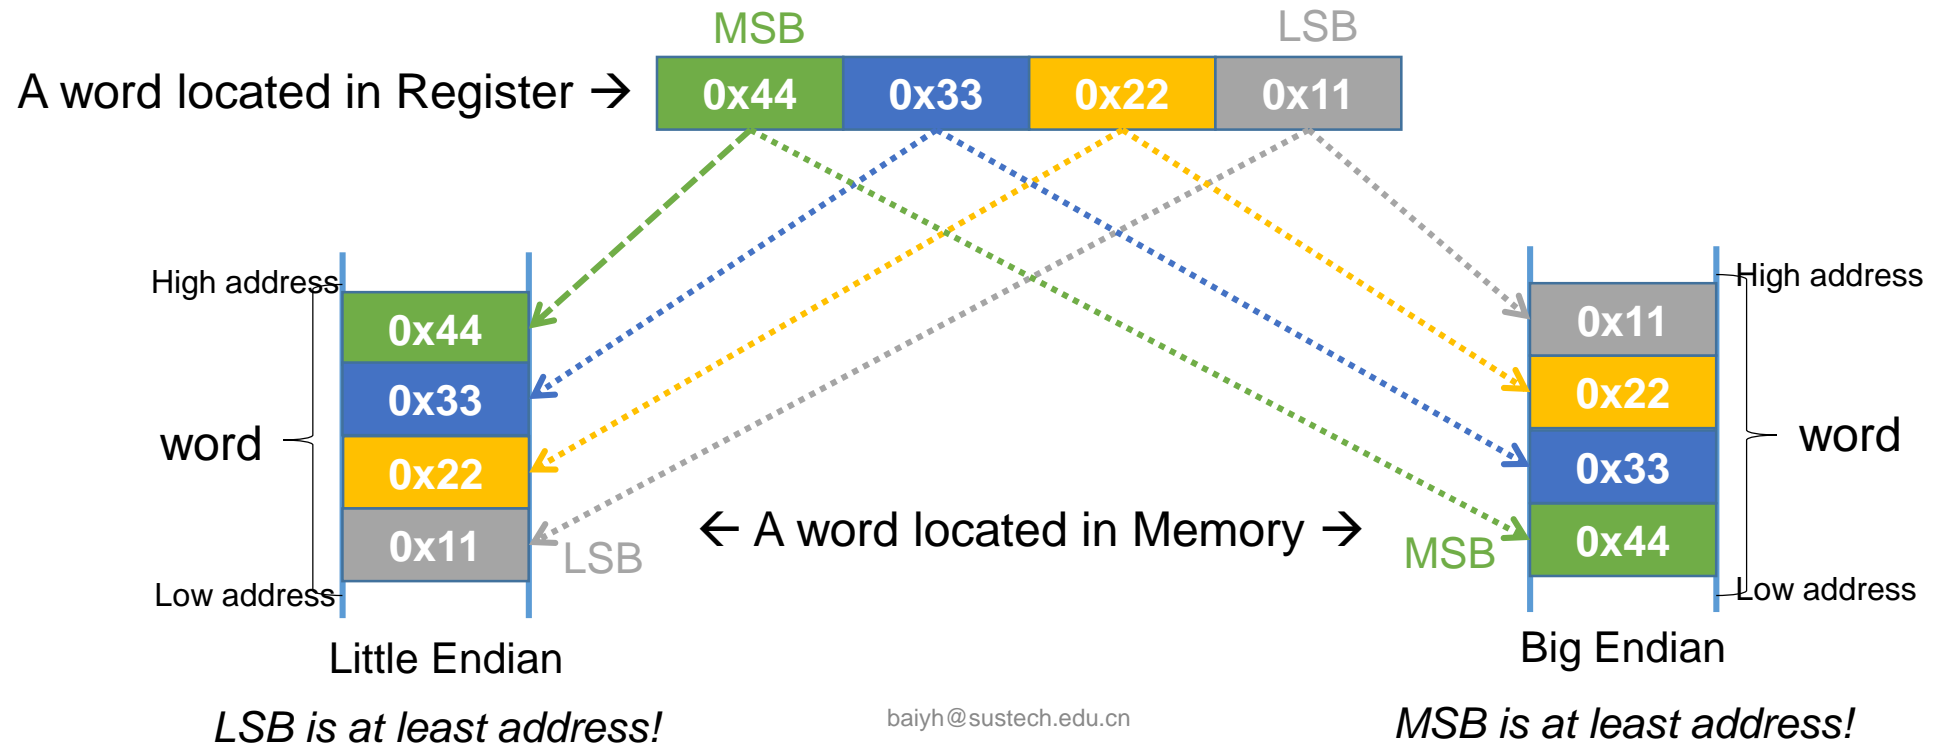
Loading and Storing Bytes

- In addition to word data transfers (lw, sw), RISC-V has byte data transfers:
 - load byte: lb
 - store byte: sb
- Same format as lw, sw
- E.g., `lb x10, 3(x11)`
 - contents of memory location (whose address = contents of register x11 + 3), is copied to the low byte position of x10.



Little Endian vs Big Endian

- Endianness: byte ordering within a word
 - Little-endian (e.g. RISC-V)
 - **LSB** of a word is at **least** memory address
 - Big-endian (e.g. MIPS)
 - **MSB** of a word is at **least** memory address



Endianness Example

- Example: For the following RISC-V code, What's the final value in x12?

```
addi x11,x0,0x3f5
```

```
sw x11,0(x5)
```

```
lb x12,1(x5)
```

- for this example

- byte[0] = 0xf5
- byte[1] = 0x03
- byte[2] = 0x00
- byte[3] = 0x00

Answer	x12
A	0x5
B	0xf
C	0x3
D	0xffffffff

Another Example

- Example: For the following RISC-V code, What's the final value in x12?

```
addi x11,x0,-512
```

```
sw x11,0(x5)
```

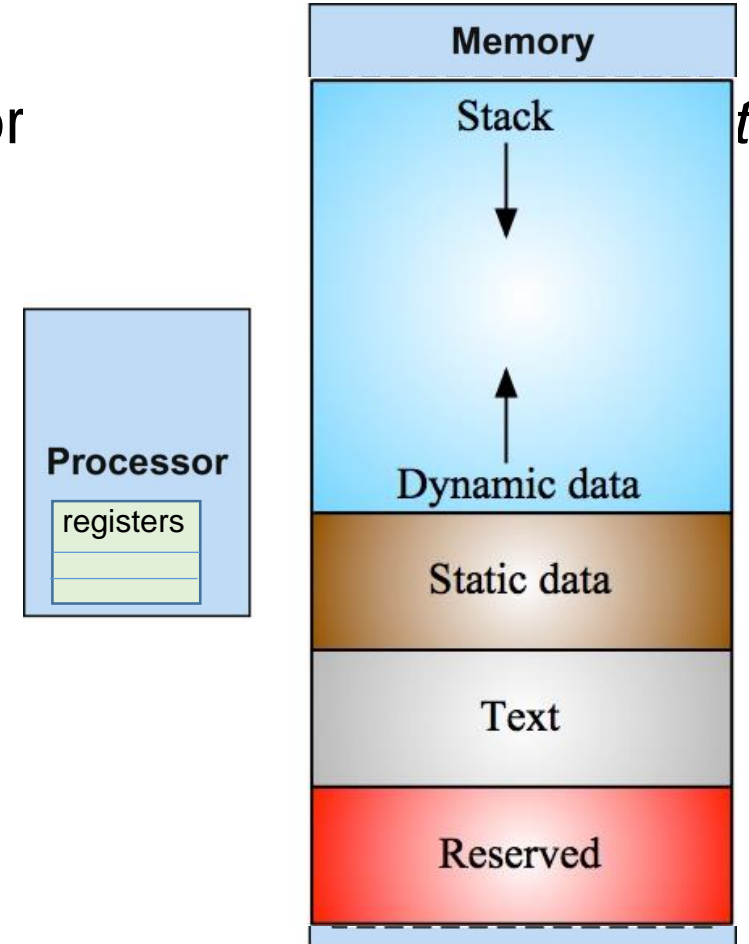
```
lb x12,1(x5)
```

- The immediate got sign extended...
 - So $-512_{10} = -2^9 = 0xffffe00$ got written
- Then load byte is called
 - So it will load byte[1], which is 0xfe
- But load byte sign extends too...
 - So what gets loaded into the register is 0xfffffffffe
- If we did lbu we'd instead get 0xfe

Answer	x12
A	0xe
B	0xfe
C	0x0
D	0xfffffffffe

Memory Layout

- Instructions(programs) are represented in binary, just like data
- Programs are stored in *Text Segment*
- Constants and other static variables are stored
- Dynamic data: *Heap*
 - E.g., malloc in C, new in Java
- Automatic data: *Stack*



Logical Operations

- Useful for extracting and inserting groups of bits in a word

Operation	C	Java	RISC-V
Shift left logical	<<	<<	sll
Shift right	>>	>>	srl/sra
Bitwise AND	&	&	and
Bitwise OR			or
Bitwise XOR	~	~	xor

- Shift left logical
 - Shift left and fill with 0 bits
 - `sll i` by i bits multiplies by 2^i
- Shift right logical
 - Shift right and fill with 0 bits
 - `srl i` by i bits divides by 2^i (unsigned only)
- Shift right arithmetic
 - Shift right and fill with sign bits

Logical vs. Arithmetic Shift

- `slli x2, x1, 4` # `reg x2 = reg x1 << 4 bits`

1001 0010 0011 0100 0101 0110 0111 1000

0010 0011 0100 0101 0110 0111 1000 0000

- `srlr x2, x1, 4` # `reg x2 = reg x1 >> 4 bits`

1001 0010 0011 0100 0101 0110 0111 1000

0000 1001 0010 0011 0100 0101 0110 0111

- `srai x2, x1, 4`

1001 0010 0011 0100 0101 0110 0111 1000

1111 1001 0010 0011 0100 0101 0110 0111

Logical Instructions

- AND: clear some bits
 - `and x9,x10,x11`
- OR: set some bits
 - `or x9,x10,x11`
- XOR: toggle some bits
 - `xor x9,x10,x12`
- How about NOT?
 - Can be implemented with XOR
 - `xori x15,x14,-1`

	0x35	0	0	1	1	0	1	0	1
AND	0x0F	0	0	0	0	1	1	1	1
	0x05	0	0	0	0	0	1	0	1

	0x04	0	0	0	0	0	1	0	0
OR	0x30	0	0	1	1	0	0	0	0
	0x34	0	0	1	1	0	1	0	0

	0x44	0	1	0	0	0	1	0	0
XOR	0x06	0	0	0	0	0	1	1	0
	0x34	0	1	0	0	0	0	1	0

Data Transfer with Variable Indexing

- C code

```
int A[100];    /* A[0] address is in x13 */
int i;         /* i in x14 */
...
g = h + A[i]; /* h = x12, g = x11, tmp = x15 */
```

- Using Load Word (lw) in RISC-V with pointer arithmetic:

```
slli x15,x14,2    # Multiply i by 4 for ints
add x15,x15,x13    # A + 4 * i
lw x10,0(x15)
add x11,x12,x10
```

Conditional Operations

- Branch to a labeled instruction if a condition is true
 - Otherwise, continue sequentially
- Conditional branch
 - *beq rs1, rs2, L1*
if (rs1 == rs2) branch to instruction labeled L1;
 - *bne rs1, rs2, L1*
 - if (rs1 != rs2) branch to instruction labeled L1;
- Unconditional branch
 - *beq x0, x0, L1*
 - unconditional jump to instruction labeled L1

Compiling If Statements


- C code

```
if (i==j) f = g+h;  
else f = g-h;
```

- i and j are in x22 and x23,
- f,g and h are in x19, x20 and x21

- Compiled RISC-V code:

```
        bne x22, x23, Else  # go to Else if i ≠ j  
        add x19, x20, x21   # f=g+h, skipped if i ≠ j  
        beq x0, x0, Exit    # unconditional go to Exit  
Else:   sub x19, x20, x21   # f=g-h, skipped if i = j  
Exit:
```



Compiling Loop Statements

- C code:

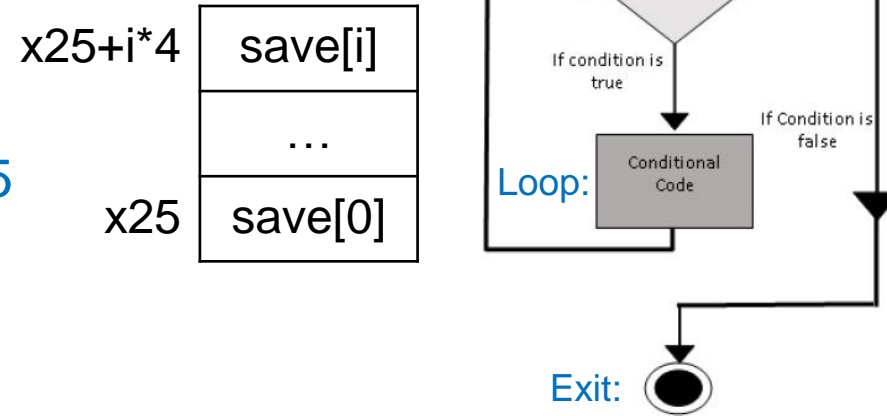
```
while (save[i] == k)
    i += 1;
```

- i in x22, k in x24, address of save in x25

- Compiled RISC-V code:

```
Loop: slli    x10, x22, 2    # Temp reg x10 = i * 4
      add     x10, x10, x25  # x10 = address of save[i]
      lw      x9, 0(x10)     # Temp reg x9 = save[i]
      bne     x9, x24, Exit  # go to Exit if save[i]≠k
      addi    x22, x22, 1    # i = i + 1
      j       Loop          # go to Loop

Exit:
```



Summary of Design Principles

- 1: Simplicity favors regularity
 - Keep all instructions the same size.
- 2: Smaller is faster
 - Register vs memory
 - Number of registers is small
- 3: Make the common case fast
 - Immediate operand
- 4: Good design demands good compromises
 - Keep formats as similar as possible

Summary

1. Instruction set architecture (ISA) specifies the set of commands (instructions) a computer can execute
2. Hardware registers provide a few very fast variables for instructions to operate on
3. RISC-V ISA requires software to break complex operations into a string of simple instructions, but enables faster, simple hardware
4. Assembly code is human-readable version of computer's native machine code, converted to binary by an assembler

RISC-V

Reference Card

- In textbook (RV64)
- Or on Blackboard (RV32)

RV32I Base Integer Instructions

Inst	Name	FMT	Opcode	funct3	funct7	Description (C)
add	ADD	R	0110011	0x0	0x00	$rd = rs1 + rs2$
sub	SUB	R	0110011	0x0	0x20	$rd = rs1 - rs2$
xor	XOR	R	0110011	0x4	0x00	$rd = rs1 \wedge rs2$
or	OR	R	0110011	0x6	0x00	$rd = rs1 rs2$
and	AND	R	0110011	0x7	0x00	$rd = rs1 \& rs2$
sll	Shift Left Logical	R	0110011	0x1	0x00	$rd = rs1 \ll rs2$
srl	Shift Right Logical	R	0110011	0x5	0x00	$rd = rs1 \gg rs2$
sra	Shift Right Arith*	R	0110011	0x5	0x20	$rd = rs1 \gg rs2$ (Arith*)
slt	Set Less Than	R	0110011	0x2	0x00	$rd = (rs1 < rs2)?1:0$
sltu	Set Less Than (U)	R	0110011	0x3	0x00	$rd = (rs1 < rs2)?1:0$
addi	ADD Immediate	I	0010011	0x0	imm[11:5]=0x00 imm[11:5]=0x00 imm[11:5]=0x20	$rd = rs1 + imm$
xori	XOR Immediate	I	0010011	0x4		$rd = rs1 \wedge imm$
ori	OR Immediate	I	0010011	0x6		$rd = rs1 imm$
andi	AND Immediate	I	0010011	0x7		$rd = rs1 \& imm$
slli	Shift Left Logical Imm	I	0010011	0x1		$rd = rs1 \ll imm[4:0]$
srli	Shift Right Logical Imm	I	0010011	0x5		$rd = rs1 \gg imm[4:0]$
srai	Shift Right Arith Imm	I	0010011	0x5		$rd = rs1 \gg imm[4:0]$ (Arith*)
slti	Set Less Than Imm	I	0010011	0x2		$rd = (rs1 < imm)?1:0$
sltiu	Set Less Than Imm (U)	I	0010011	0x3		$rd = (rs1 < imm)?1:0$
lb	Load Byte	I	0000011	0x0		$rd = \{24'bM[rs1+imm][7], M[rs1+imm][7:0]\}$
lh	Load Half	I	0000011	0x1		$rd = \{16'bM[rs1+imm][15], M[rs1+imm][15:0]\}$
lw	Load Word	I	0000011	0x2		$rd = M[rs1+imm][31:0]$
lbu	Load Byte (U)	I	0000011	0x4		$rd = \{24'b0, M[rs1+imm][7:0]\}$
lhu	Load Half (U)	I	0000011	0x5		$rd = \{16'b0, M[rs1+imm][15:0]\}$
sb	Store Byte	S	0100011	0x0		$M[rs1+imm][7:0] = rs2[7:0]$
sh	Store Half	S	0100011	0x1		$M[rs1+imm][15:0] = rs2[15:0]$
sw	Store Word	S	0100011	0x2		$M[rs1+imm][31:0] = rs2[31:0]$
beq	Branch ==	B	1100011	0x0		$\text{if}(rs1 == rs2) \text{ PC} = \text{PC} + \{imm, 1'b0\}$
bne	Branch !=	B	1100011	0x1		$\text{if}(rs1 != rs2) \text{ PC} = \text{PC} + \{imm, 1'b0\}$
blt	Branch <	B	1100011	0x4		$\text{if}(rs1 < rs2) \text{ PC} = \text{PC} + \{imm, 1'b0\}$
bge	Branch ≥	B	1100011	0x5		$\text{if}(rs1 \geq rs2) \text{ PC} = \text{PC} + \{imm, 1'b0\}$
bltu	Branch < (U)	B	1100011	0x6		$\text{if}(rs1 < rs2) \text{ PC} = \text{PC} + \{imm, 1'b0\}$
bgeu	Branch ≥ (U)	B	1100011	0x7		$\text{if}(rs1 \geq rs2) \text{ PC} = \text{PC} + \{imm, 1'b0\}$
jal	Jump And Link	J	1101111			$rd = \text{PC}+4; \text{ PC} = \text{PC} + \{imm, 1'b0\}$
jalr	Jump And Link Reg	I	1101111	0x0		$rd = \text{PC}+4; \text{ PC} = rs1 + imm$
lui	Load Upper Imm	U	0110111			$rd = imm \ll 12$
auipc	Add Upper Imm to PC	U	0010111			$rd = \text{PC} + (imm \ll 12)$
ecall	Environment Call	I	1110011	0x0	imm=0x0	Transfer control to OS
ebreak	Environment Break	I	1110011	0x0	imm=0x1	Transfer control to debugger